




Generative adversarial networks

22-808: Generative models
Sharif University of Technology
Fall 2024

Fatemeh Seyyedsalehi

Recap



- ▶ We need a framework to interact with distributions for statistical generative models.
 - ▶ Probabilistic generative models
 - ▶ Deep generative models
 - ▶ Autoregressive models $p_{\theta}(\mathbf{x}) = \prod_{i=1}^n p_{\theta}(x_i | \mathbf{x}_{<i})$
 - ▶ Variational Autoencoders $p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{x}, \mathbf{z}) d\mathbf{z}$
 - ▶ **Generative adversarial networks**
 - ▶ Both AR and VAE model families attempted to minimize the KL divergence between model family and data distribution, or equivalently attempt to maximize the likelihood.
 - ▶ In GAN we are going to use an alternative choice for the similarity measure between model distribution and data distribution.


Maximizing the likelihood

$$\hat{\theta} = \operatorname{argmax}_{\theta} \sum_{i=1}^M \log p_{\theta}(\mathbf{x}_i), \quad \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_M \sim p_{\text{data}}(\mathbf{x})$$

- ▶ Optimal statistical efficiency
 - ▶ Assume sufficient model capacity, such that there exists a unique θ^* that satisfy $p_{\theta^*} = p_{\text{data}}$.
 - ▶ The convergence of $\hat{\theta}$ to θ^* when $M \rightarrow \infty$, is the fastest among all statistical methods when using maximum likelihood training.

Maximizing the likelihood


- ▶ For imperfect models, achieving high log-likelihoods might not always imply good sample quality.



An isotropic Gaussian distribution was fit to data drawn from a mixture of Gaussians by either minimizing KL divergence (KLD), maximum mean discrepancy (MMD), or Jensen-Shannon divergence (JSD). The different fits demonstrate different tradeoffs made by the three measures of distance between distributions.

Implicit generative models

- ▶ Kind of probabilistic generative models without an explicit likelihood function
- ▶ We use a likelihood-free approach to train these models
 - ▶ Training by comparing samples



Explicit models vs. implicit models

Learning by comparing samples

- ▶ We should define a distance(similarity) measure between two distributions that:
 - ▶ Provides guarantees about learning the data distribution.

$$\operatorname{argmin}_{p_{\theta}} D(p_{data}, p_{\theta}) = p_{data}$$

- ▶ Can be evaluated only using samples from the data and model distribution.
 - ▶ Are computationally cheap to evaluate.
- ▶ Many distributional distances and divergences fail to satisfy the later two requirements

Learning by comparing samples

- ▶ The main approach to overcome these challenges is to approximate the desired quantity through optimization by introducing a comparison model, often called a **discriminator** or a **critic** D , such that:

$$\mathcal{D}(p^*, q) = \operatorname{argmax}_D \mathcal{F}(D, p^*, q)$$

- ▶ where \mathcal{F} is a functional that can be estimated using only samples from $p^*(p_{data})$ and q . One way is that it depends on distributions only in expectations.
 - ▶ Therefore, it can be estimated using Monte Carlo estimation.

Learning by comparing samples

- ▶ As we usually use parametric functions (ex. Neural networks) for both the model and discriminator.
- ▶ Therefore, by the following optimization we estimate the distance measure $\mathcal{D}(p^*, q_\theta)$


$$\operatorname{argmax}_{\phi} \mathcal{F}(D_{\phi}, p^*, q_{\theta})$$

- ▶ Then, instead of optimizing the exact objective $\mathcal{D}(p^*, q_\theta)$ we use the tractable approximation provided through the optimal D_{ϕ} .

Generative adversarial networks

(Goodfellow GAN)

- ▶ A finite number of samples from the desired real distribution is available: x_1, x_2, \dots, x_n
- ▶ Like VAEs, we consider a latent variable model for the model generation process and attempt to learn G_θ . However, here we learn this function by Comparing samples.



The Goodfellow GAN

The probabilistic classification view

- ▶ Assuming $D(x)$ as a binary classifier which predicts whether a given point x was sampled from the real distribution or it is a fake sample from the generator G_θ .

- ▶ A cross entropy loss to train this classifier:

$$E_{\mathbf{x} \sim p_{\text{data}}} [\log D_\phi(\mathbf{x})] + E_{\mathbf{x} \sim p_\theta} [\log(1 - D_\phi(\mathbf{x}))]$$

- ▶ We can see that the optimal discriminator for a fixed generator G_θ is:

$$\frac{p(x)}{p(x) + p_\theta(x)}$$

The Goodfellow GAN

The objective function

- By substitution the optimal discriminator into the cross-entropy loss, we have:

$$\begin{aligned} V^*(q_\theta, p^*) &= \frac{1}{2} \mathbb{E}_{p^*(\mathbf{x})} \left[\log \frac{p^*(\mathbf{x})}{p^*(\mathbf{x}) + q_\theta(\mathbf{x})} \right] + \frac{1}{2} \mathbb{E}_{q_\theta(\mathbf{x})} \left[\log \left(1 - \frac{p^*(\mathbf{x})}{p^*(\mathbf{x}) + q_\theta(\mathbf{x})} \right) \right] \\ &= \frac{1}{2} \mathbb{E}_{p^*(\mathbf{x})} \left[\log \frac{p^*(\mathbf{x})}{\frac{p^*(\mathbf{x}) + q_\theta(\mathbf{x})}{2}} \right] + \frac{1}{2} \mathbb{E}_{q_\theta(\mathbf{x})} \left[\log \left(\frac{q_\theta(\mathbf{x})}{\frac{p^*(\mathbf{x}) + q_\theta(\mathbf{x})}{2}} \right) \right] - \log 2 \\ &= \frac{1}{2} D_{\text{KL}} \left(p^* \parallel \frac{p^* + q_\theta}{2} \right) + \frac{1}{2} D_{\text{KL}} \left(q_\theta \parallel \frac{p^* + q_\theta}{2} \right) - \log 2 \\ &= JSD(p^*, q_\theta) - \log 2 \end{aligned}$$

where JSD is the Jensen-Shannon divergence.

The Goodfellow GAN

- ▶ This establishes a connection between optimal binary classification and distributional divergences.
- ▶ By using binary classification, we were able to compute the distributional divergence using only samples, which is the important property needed for learning implicit generative models
- ▶ We have turned an intractable estimation problem (how to estimate the JSD divergence) into an optimization problem (how to learn a classifier) which can be used to approximate that divergence.

The Goodfellow GAN

- ▶ With optimal discriminator, we attempt to find the generative model G_θ that minimizes the JSD divergence.

$$\begin{aligned}\min_{\theta} JSD(p^*, q_\theta) &= \min_{\theta} V^*(q_\theta, p^*) + \log 2 \\ &= \min_{\theta} \frac{1}{2} \mathbb{E}_{p^*(\mathbf{x})} \log D^*(\mathbf{x}) + \frac{1}{2} \mathbb{E}_{q_\theta(\mathbf{x})} \log(1 - D^*(\mathbf{x})) + \log 2\end{aligned}$$

Training procedure of GAN

Sample minibatch of m training points $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}$ from \mathcal{D}

Sample minibatch of m noise vectors $\mathbf{z}^{(1)}, \mathbf{z}^{(2)}, \dots, \mathbf{z}^{(m)}$ from p_z

Update the discriminator parameters ϕ by stochastic gradient **ascent**

$$\nabla_{\phi} V(G_{\theta}, D_{\phi}) = \frac{1}{m} \nabla_{\phi} \sum_{i=1}^m [\log D_{\phi}(\mathbf{x}^{(i)}) + \log(1 - D_{\phi}(G_{\theta}(\mathbf{z}^{(i)})))]$$

Update the generator parameters θ by stochastic gradient **descent**

$$\nabla_{\theta} V(G_{\theta}, D_{\phi}) = \frac{1}{m} \nabla_{\theta} \sum_{i=1}^m \log(1 - D_{\phi}(G_{\theta}(\mathbf{z}^{(i)})))$$

Repeat for fixed number of epochs

Active

Training convergence

- ▶ If G and D have enough capacity, and at each step of training procedure, the discriminator is allowed to reach its optimum for a specific G_θ , and then p_θ is updated so as to improve


$$\mathbb{E}_{\mathbf{x} \sim p_{data}} [\log D_G^*(\mathbf{x})] + \mathbb{E}_{\mathbf{x} \sim p_g} [\log(1 - D_G^*(\mathbf{x}))]$$

then p_θ converges to p_{data} .

- ▶ Unrealistic assumptions ☹️

Training convergence

- ▶ However, we do not have access to the optimal discriminator and only we can approximate it with a parametrized function: neural network D_ϕ
- ▶ No guarantee for convergence
- ▶ In practice, the generator and discriminator loss keeps oscillating during GAN training




The min-max game

► The minmax game

$$\min_{\theta} \max_{\phi} V(G_{\theta}, D_{\phi}) = E_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + E_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

- It is a game not an optimization problem
- It should reach to a Nash equilibria



Example

- ▶ Which one is real?




F-divergence

- ▶ Let $f: \mathbb{R} \rightarrow \mathbb{R}$ be a convex lower-semicontinuous function, such that $f(1) = 0$. We define the *f-divergence* between two distributions with densities p and q by:

$$D_f(p \parallel q) \equiv \int_{\mathcal{X}} q(x) f\left(\frac{p(x)}{q(x)}\right) dx.$$

- ▶ What's interesting about *f-divergence* is that we can construct a variational representation for it.
 - ▶ Alternating the integral to an optimization

Convex lower-semicontinuous function



Fenchel duality

- ▶ The idea is to use the convex conjugate of the function f , which is defined as follows:

$$f^*(t) \equiv \sup_x \{tx - f(x)\}.$$

- ▶ Fenchel duality: repeat application of the conjugate operation to convex lower-semicontinuous function f yields $f^{**} = f$. Therefore, we have:

$$f(x) = \sup_t \{tx - f^*(t)\}.$$

Variational representation of F -divergence

- ▶ Using Fenchel duality, we obtain the variational representation of the f -divergence.

$$\begin{aligned} D_f(p \parallel q) &= \int_{\mathcal{X}} q(x) \sup_t \left[t \frac{p(x)}{q(x)} - f^*(t) \right] dx \\ &= \int_{\mathcal{X}} \sup_t [tp(x) - f^*(t)q(x)] dx \\ &= \sup_{T: \mathcal{X} \rightarrow \mathbb{R}} \int_{\mathcal{X}} (T(x)p(x) - f^*(T(x))q(x)) dx \\ &= \sup_{T: \mathcal{X} \rightarrow \mathbb{R}} \left[\mathbb{E}_{x \sim p} T(x) - \mathbb{E}_{x \sim q} f^*(T(x)) \right]. \end{aligned}$$

F-GAN

- ▶ The dual form can be approximated using Monte Carlo estimation.
- ▶ Assuming a parametric family of functions $T\varphi$ (ex. a neural network) and the generator function g_θ , and a valid f-divergence, the F-GAN objective is,

$$\begin{aligned}\theta_f &= \arg \min_{\theta} \sup_{\varphi} \left[\mathbb{E}_{x \sim p} T_{\varphi}(x) - \mathbb{E}_{x \sim p_{\theta}} f^*(T_{\varphi}(x)) \right] \\ &= \arg \min_{\theta} \sup_{\varphi} \left[\mathbb{E}_{x \sim p} T_{\varphi}(x) - \mathbb{E}_{z \sim q} f^*(T_{\varphi}(g_{\theta}(z))) \right].\end{aligned}$$

- ▶ Generator g_θ tries to minimize the divergence estimate and discriminator $T\varphi$ tries to tighten the lower bound

F-divergence

distance or divergence	corresponding $g(t)$ ($t = \frac{p_i(x)}{p_j(x)}$)
Bhattacharyya distance ¹	\sqrt{t}
KL-divergence	$t \log(t)$
Symmetric KL-divergence	$t \log(t) - \log(t)$
Hellinger distance	$(\sqrt{t} - 1)^2$
Total variation	$ t - 1 $
Pearson divergence	$(t - 1)^2$
Jensen-Shannon divergence	$\frac{1}{2} (t \log \frac{2t}{t+1} + \log \frac{2}{t+1})$

The Goodfellow GAN as F-GAN

- ▶ The Goodfellow GAN is an instances of the f -GAN.
- ▶ Modified version of the Jensen-Shannon

$$2\text{JSD}(p, q) - \log(4) = D_{\text{KL}} \left(p \left\| \frac{p+q}{2} \right. \right) + D_{\text{KL}} \left(p_g \left\| \frac{p+q}{2} \right. \right) - \log(4).$$

- ▶ The f -divergence:

$$f(x) = x \log x - (x+1) \log(x+1)$$

$$f^*(t) = -\log(1 - e^t).$$


$$T_\varphi(x) = \log(d_\varphi(x))$$

- ▶ We can obtain the Goodfellow GAN :


$$_{25} \theta_f = \arg \min_{\theta} \sup_{\varphi} \left[\mathbb{E}_{x \sim p} \log d_\varphi(x) + \mathbb{E}_{z \sim q} \log(1 - d_\varphi(g_\theta(z))) \right]$$

Mode collapse and catastrophic forgetting


- ▶ In the case of mode collapse, the generator might focus on producing only a limited set of outputs that it knows will deceive the discriminator, completely ignoring other parts of the data distribution.
- ▶ As the generator iterates over epochs, it starts to forget the diversity it initially captured.
 - ▶ This happens because it gets reinforced to produce only certain types of outputs that are effective in fooling the discriminator.




Target




Step 0



Step 5k



Step 10k



Step 15k




Step 20k

Source: Metz et al., 2017

The problem with KL divergence


- ▶ KL divergence problem:
 - ▶ When distributions' supports are different, the KL does not defined.
 - ▶ As it consider the ratio of probability values, it shows a big difference between two distributions when one has a very small value in even in a small region.




Wasserstein GAN

- ▶ Earth-Mover (EM) distance (Wasserstein-1)
 - ▶ $\Pi(P_r, P_g)$ shows the set of all joint distributions whose marginals are P_r and P_g , respectively.

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$



- ▶ It is the cost of optimal transport between two distributions P_r and P_g .



Jeremy Kun, (blog post, 2018)

Example

- ▶ Consider $z \sim U[0,1]$
 - ▶ P_0 a distribution over $(0, z)$
 - ▶ P_θ a distribution over (θ, z)
- ▶ Different distance measure for these two distributions:

- $W(\mathbb{P}_0, \mathbb{P}_\theta) = |\theta|,$



- $JS(\mathbb{P}_0, \mathbb{P}_\theta) = \begin{cases} \log 2 & \text{if } \theta \neq 0, \\ 0 & \text{if } \theta = 0, \end{cases}$

- $KL(\mathbb{P}_\theta \| \mathbb{P}_0) = KL(\mathbb{P}_0 \| \mathbb{P}_\theta) = \begin{cases} +\infty & \text{if } \theta \neq 0, \\ 0 & \text{if } \theta = 0, \end{cases}$

- and $\delta(\mathbb{P}_0, \mathbb{P}_\theta) = \begin{cases} 1 & \text{if } \theta \neq 0, \\ 0 & \text{if } \theta = 0. \end{cases}$

Example

- ▶ Learning can not be done with the other distances and divergences because the resulting loss function is not even continuous.
- ▶ Comparing EM and JSD for different θ



Kantorovich-Rubinstein Duality

- ▶ We can approximate the Wasserstein distance with its dual form:

$$W(\mathbb{P}_r, \mathbb{P}_\theta) = \sup_{\|f\|_L \leq 1} \mathbb{E}_{x \sim \mathbb{P}_r}[f(x)] - \mathbb{E}_{x \sim \mathbb{P}_\theta}[f(x)]$$


where the sup is over all 1-Lipschitz functions $f : X \rightarrow \mathbb{R}$.

- ▶ Considering a parameterized family of functions for f :
 - ▶ However, we need to be sure that this family satisfy the 1-Lipschitz constraint.
 - ▶ The Lipschitz constraint is essentially that a function must have a maximum gradient. The specific maximum gradient is a hyperparameter.

$$\max_{w \in \mathcal{W}} \mathbb{E}_{x \sim \mathbb{P}_r}[f_w(x)] - \mathbb{E}_{z \sim p(z)}[f_w(g_\theta(z))]$$

Constraint on the discriminator function

- ▶ The red line is a good discriminator but its gradient is nearly 0 at most points. The cyan line is clearly much worse as a discriminator, but is much better for training the generator because its gradient is not zero.
- ▶ The Lipschitz constraint limits the discriminator function



Lipschitz constraint

- ▶ Quick and dirty solution: clamp the size of the weights

$$-c < W < c$$

- ▶ Or clipping the gradient.
- ▶ However, a better solution is to add a soft penalty to the loss function as follows: (WGAN-GP)

$$\underbrace{\mathbb{E}_{\tilde{\mathbf{x}} \sim \mathbb{P}_g} [D(\tilde{\mathbf{x}})] - \mathbb{E}_{\mathbf{x} \sim \mathbb{P}_r} [D(\mathbf{x})]}_{\text{Original critic loss}} + \lambda \underbrace{\mathbb{E}_{\hat{\mathbf{x}} \sim \mathbb{P}_{\hat{\mathbf{x}}}} [(\|\nabla_{\hat{\mathbf{x}}} D(\hat{\mathbf{x}})\|_2 - 1)^2]}_{\text{Our gradient penalty}}$$

Where $\hat{\mathbf{x}}$ is uniformly sampled from the line between samples of two distributions